Toy Models of Superposition Replication and Findings

Zephaniah Roe   zroe@uchicago.edu

Undergraduate Student at the University of Chicago

Abstract

Toy Models of Superposition[1] is a groundbreaking paper published by researchers affiliated with Anthropic and Harvard in 2022. By investigating small models with under 100 neurons, the paper demonstrates that neural networks can represent more features than they have dimensions. Additionally, they use these so-called "toy models" to understand the relationship between how neural networks are trained and how they represent the data internally. Because the paper is quite extensive, this replication only focuses on reproducing the most important results from the introduction and sections 2 and 3 of the original paper. It also includes some commentary on section 1.

1   Introduction

Toy Models of Superposition motivates the idea of superposition with the following graphic:
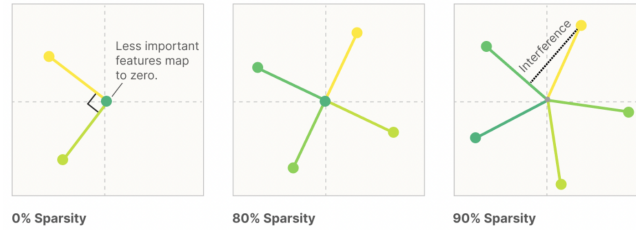

Figure 1: Graphic from *Toy Models of Superposition* showing superposition in 2D.

The basic idea is this: if you think of each feature as being represented inside of a neural network by a direction, you can graph these directions and observe them. By doing this, the authors of *Toy Models of Superposition* demonstrate that the internal structure of a model depends on the sparsity of its training data. A replication of this phenomenon can be found below and the code used to generate it can be found here.
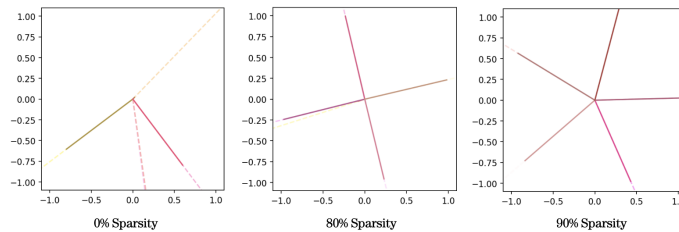

Figure 2: Replication of figure 1

The model studied in figure 2 is designed such that each column in the weight matrix corresponds to a given input. The weight matrix in this example contains only 2 neurons so each column can be graphed as a 2D vector. Observing these vectors while increasing the sparsity of the model's input reveals that the model can represent five features despite having only 2 neurons. This kind of representation is what the authors call "superposition." In future sections, we will study superposition extensively and reproduce this phenomenon in larger networks.

2 Background and Motivation

Before studying superposition in detail, the authors of *Toy Models of Superposition*, provide context by explaining concepts and defining terms. In this section, I provide some additional commentary on the definitions and explanations the authors use.

(1) Defining Features: *Toy Models of Superposition* defines features broadly as "properties of the input which a sufficiently large neural network will reliably dedicate a neuron to representing." The authors do, however, describe this definition as "slightly circular" and note that they are not "overly attached to it." I find the definition especially problematic because certain architectures may be incentivized to ignore an aspect of the input that differently designed models may "want" to internally represent. As a result, I propose an alternative definition: features are aspects of the input that a neural network represents accurately with a significantly higher probability than a randomly initialized network. In other words, features are parts of the input that a model determines to be important enough to represent internally. (This definition is very similar to the "features as arbitrary functions" definition from the original paper.)

(2) Role of Linear Representations in Neural Networks: The original authors of the paper study interpretability by trying to understand the linear representations within neural networks. It is worth noting that this isn't the only way to approach mechanistic interpretability research. Understanding the role of non-linearities at each level is likely also very important (and perhaps more neglected).

(3) Defining Superposition: The original paper has a compelling yet simple definition for superposition: "Roughly, the idea of superposition is that neural networks 'want to represent more features than they have neurons', so they exploit a property of high-dimensional spaces to simulate a model with many more neurons." This is the definition I will use throughout this paper.

3 Demonstrating Superposition

In the introduction, the authors of the original paper proved that models with two neurons could exhibit superposition (this result was reproduced in figure 2). Later in the paper, however, the authors demonstrate that superposition is also observed in models with more than two neurons. Specifically, they begin by studying a linear model defined by $W^T W x + b$ and a ReLU model defined by $\text{ReLU}(W^T W x + b)$—both with a weight matrix $W_{5 \times 20}$. The objective of both models was to reconstruct the input x. The authors used a weighted mean squared error loss function making accurately representing some features more important than others. For more information about the loss see section 2 of the original paper.

While the model without an activation function appears to only represent features orthogonally, the authors claim that the ReLU model can exhibit superposition if it is trained on sparse enough data. Because we defined each weight matrix with 5 neurons ($W_{5 \times 20}$) this claim cannot be validated by graphing features in 2D like in figure 2. But by graphing $W^T W$, superposition in the ReLU model can be shown visually in figure 3. In this figure, positive numbers in the matrix $W^T W$ are labeled red while negative ones are colored blue. The authors also graph the length of each feature by treating each column in W as a vector. Features that are orthogonal to others in W are labeled black while features that aren't

are labeled yellow (the exact details for how this is calculated is discussed in 3.1).
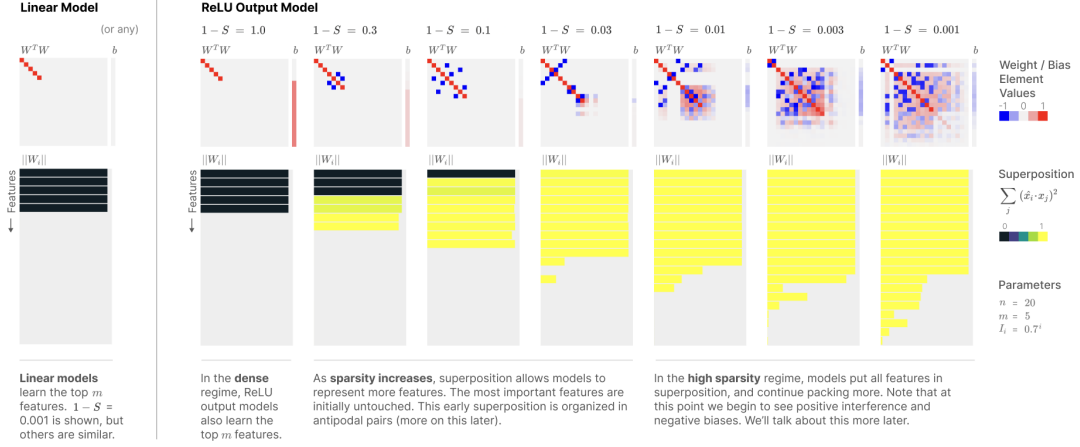

Figure 3: Superposition in linear and ReLU models from *Toy Models of Superposition*[1]

The visualization of $W^T W$ (top of figure 3) and the chart of feature representations (bottom of the figure) both show that by increasing sparsity, the ReLU model ceases to represent features orthogonally. The yellow bars shown in the models on the right side of the figure illustrate that the model maps all its features in superposition. This is also shown in the graphs of the weight matrices $W^T W$ for sparse models in figure 3. These representations appear to show the model embedding more than 5 features, but the representation is noisy.

The first step in replicating these findings was to train the linear and ReLU models that don't perform computation in superposition. The objective of each model was to reconstruct the input x. Both models were trained with the Adam optimizer (learning rate $= 1 * 10^{-3}$) on 20,000 batches of 256 examples.
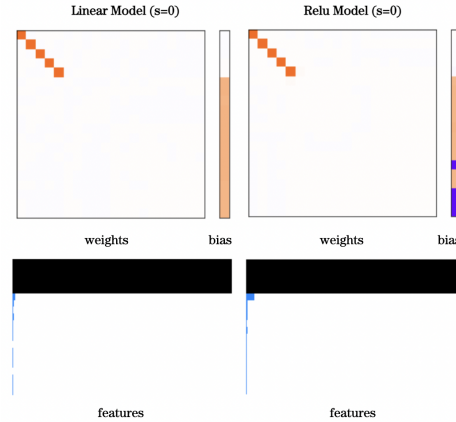

Figure 4: The generated graphics show that the linear and ReLU models trained on dense inputs represent only five features (one for each dimension in the model).

Note that in figure 4, I use orange to indicate positive numbers and purple to indicate negative ones. This is different from the red and blue in figure 3 to distinguish my work from that of the original authors. Similarly, while the original authors use yellow to indicate features in superposition, I use blue (This is hard to see in figure 4 but it will be more obvious going forward).

3.1 Calculating Superposition

The models in figure 4 do not exhibit superposition. They encode the five most important features orthogonally (one feature for each neuron in the model). In this section, we will be investigating models that do not behave in this way, instead encoding features as vectors that interfere with each other. The model explored in this section is defined by the same architecture and objective as the models in figure 4. The difference is that the models in this section are trained on sparse input data and, as a result, map features in superposition.

In order to explain this phenomenon and demonstrate how models with sparse input are able to represent features in superposition, it will be useful to dive into the math behind the concept of feature interference. The extent to which features interfere with each other is defined by the following equation:

$$\text{Interference} = \sum_{j \neq i} (\hat{W}_i \cdot \hat{W}_j)^2 \tag{1}$$

For a given column i in weight matrix W, interference is calculated by taking the dot product with every other column in W. Non-zero dot products indicate that the columns in W are not orthogonal. As a result, summing these dot products gives a general idea of how much the network is representing a given feature in superposition. Note that $\hat{W}_i$ is the unit vector for W_i. This is necessary because when calculating interference, we are interested in the direction of a given feature, not its length.

In figure 5, the length of a feature (calculated by taking the length of the vector W_i) determines the width of the bars in the feature graph (shown in the bottom half of the figure). The interference equation (Equation 1) determines the color of the columns: black indicates a low value for interference while blue indicates a higher value. This means that blue bars show that a given feature is represented in superposition while black bars indicate that the feature is mapped orthogonally.
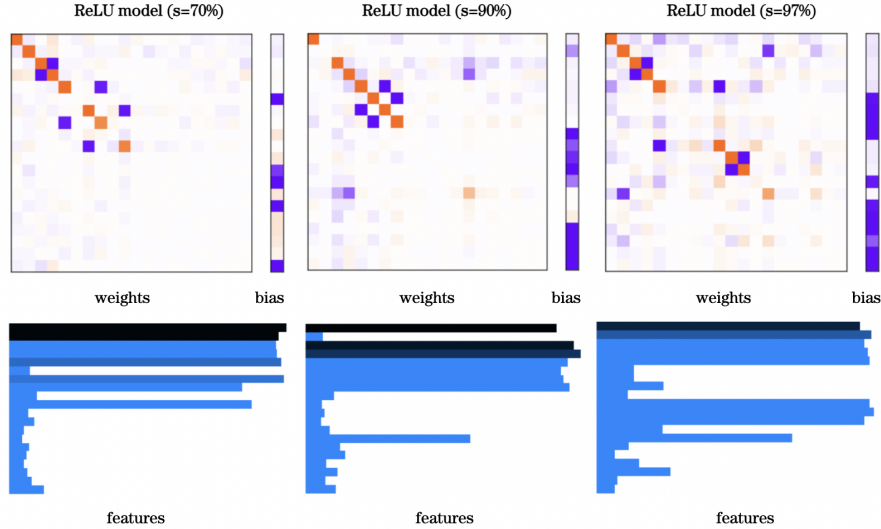

Figure 5: Superposition is observed in models trained on 70%, 90% and 97% sparse inputs (code to generate these figures can be found here). The 70% and 90% sparse models were trained on 50,000 batches of 256 examples with the "RMSProp" optimizer (learning rate $= 10^{-2}$). The 97% sparsity model was trained on 100,000 batches of 256 examples using the Adam optimizer (learning rate $= 10^{-2}$).

Unlike the models with 0% sparsity in figure 4, the models in figure 5 have higher levels of sparsity and, as a result, leverage superposition. The bottom half of figure 5 shows that these models represent far

4

more features than the models in figure 4, but by doing so, they are forced to represent many of their features in superposition. The models only have 5 neurons so if they "want" to represent more than 5 features, they can't represent each feature orthogonally. This tradeoff is intuitively more attractive when the model is trained on sparse inputs because it is less likely that the model will be fed a combination of inputs that cause feature representations to conflict (because a significant percentage of the input is 0).

3.2 Models Trained on Very Sparse Data

As sparsity is increased to almost 100% the models stop representing any features orthogonally. This is displayed in figure 6 where models are trained on 99%, 99.7% and 99.9% sparse inputs.
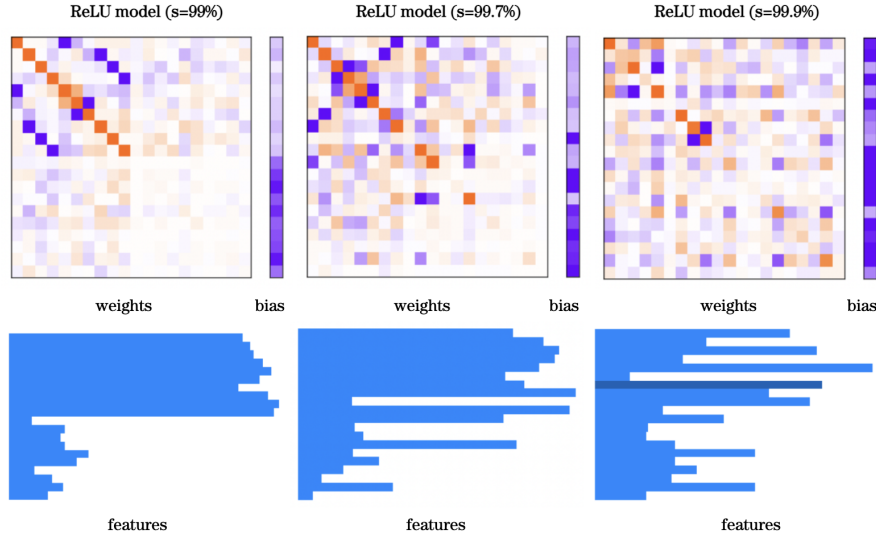

Figure 6: When models are trained on sufficiently sparse data, all feature representations are in superposition. The models in the figure were trained on 100,000 batches of 256 examples using the Adam optimizer (learning rate = 10^{-2}).

The representations of $W^T W$ of these very sparse models (shown in the top half of figure 6) is far less clean than previous representations we have seen. This is the same trend the original authors found when increasing sparsity of these models (figure 3 illustrates how the original authors displayed this visually).

The feature representations, shown in the bottom half of figure 6, are also consistent with the findings of *Toy Models of Superposition*. Like the investigation from the original paper, these feature representations show no features mapped orthogonally (recall that features that interfere with each other are shown in blue).

3.3 Scaling Results to Larger Models

In the previous subsections, we have explored superposition in models with 5 neurons and 20 inputs. This proves that the phenomenon we observed in figure 2 applies to models with more than 2 neurons. The original paper expanded this finding by also studying models with 20 neurons and 80 inputs (figure 7).
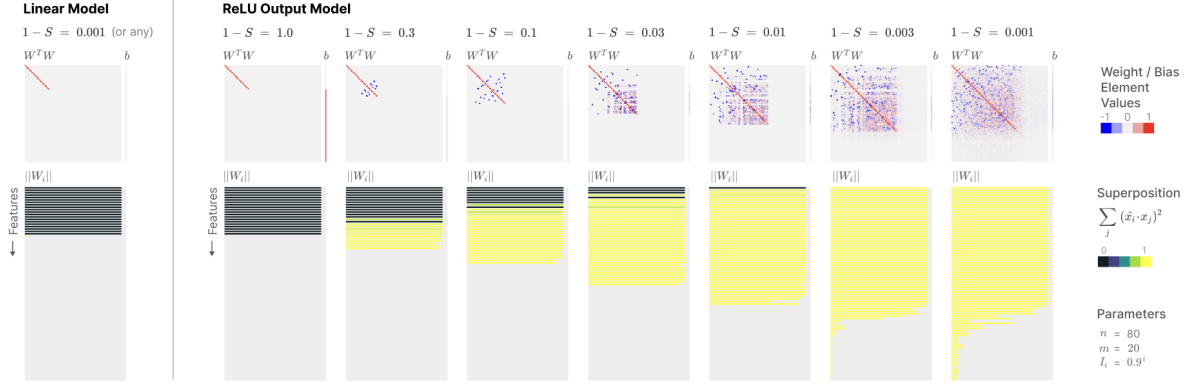
Figure 7: Superposition in models with 20 neurons and 80 inputs from *Toy Models of Superposition*[1]

The original authors report that this experiment produced "qualitatively similar" results compared to the models with 5 neurons and 20 inputs. Because this experiment is almost identical to the previous ones, and because I would not expect to find results that conflict with the original paper, I have excluded this experiment from this replication to focus on other sections.

4   Superposition as a Phase Change

Important: **I am still developing section 4 of this replication. Subsections 4, 4.1, 4.2, 4.3 and 4.4 represent a complete draft of my initial thoughts and findings. My beliefs about some of the claims in these sections have already been updated but have not yet been included in this document.**

The author's of *Toy Models of Superposition* claim that superposition can be thought of as a kind of "phase change." Figure 8 shows a phase diagram for single-neuron models with 2 inputs. The models studied were defined by $\text{ReLU}(W^TWx+b)$ and trained to simply reconstruct their inputs. The authors used a weighted mean squared error loss where the importance of the second output was varied between 0.1 and 10. The relative importance of this second output represents the x-axis for the graphs in figure 8 while he y-axis reflects the model's feature density—in other words, 1 - sparsity.
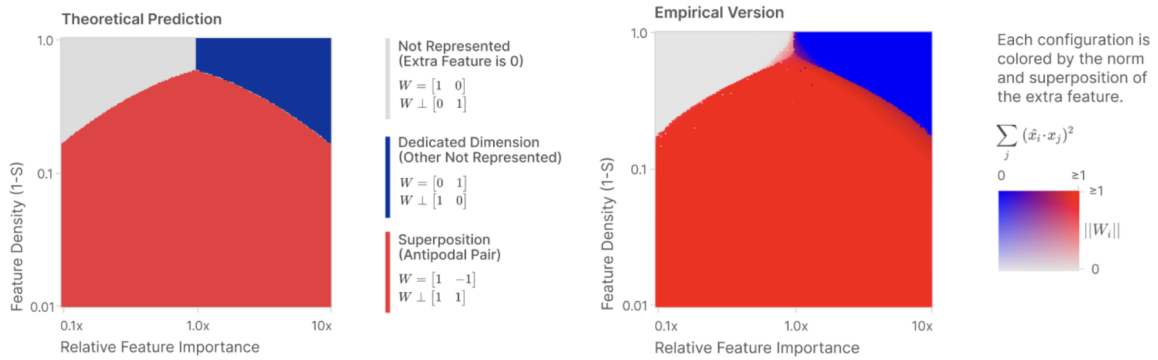

Figure 8: Phase diagrams for one-neuron models in *Toy Models of Superposition*[1]

These phase diagrams present three discrete possibilities for the way this single neuron model represents

6

it's input internally. The colors on the graph represent which of these internal configurations lead to the lowest theoretical loss. The first possibility is the model *only* represents the first input. In this case, the weight matrix $W = [1, 0]$ (shown in gray in figure 8). The opposite option is also a possibility: the model could only embed the second feature making $W = [0, 1]$ (shown in blue). The third possibility is the model could represent *both* features by making $W = [1, -1]$ (shown in red).

The authors claim it is possible to calculate theoretical losses based on sparsity, relative feature importance, and the weight matrix W for each of these models. *Toy Models of Superposition* links to this notebook which specifies the equations for calculating these theoretical losses for models with 3 inputs and 2 neurons, but they should work exactly the same in this simpler example:

$$\text{Loss for when } W \text{ is } [1, 0] = \frac{s}{3} - \frac{s^2}{4} \tag{2}$$

$$\text{Loss for when } W \text{ is } [0, 1] = r\left(\frac{s}{3} - \frac{s^2}{4}\right) \tag{3}$$

$$\text{Loss for when } W \text{ is } [1, -1] = \frac{(1+r)s^2}{6} \tag{4}$$

The variable r is the relative importance of the second feature. The variable s is the feature density—in other words, 1 - sparsity. For information about how these equations were derived, visit the notebook provided by the authors. Based on these equations, the authors were able to plot which weight configuration would theoretically lead to the lowest loss in figure 8. They were then able to train models and average the results to show that the same pattern emerges empirically when training models with gradient descent.

4.1 Replicating Superposition as a Phase Change

Replicating the findings from *Toy Models of Superposition* in the previous subsection was an intricate process that will be described in detail in the succeeding sections. Figure 9 is a visual summary of my findings. It will be referred to in the subsections below as a way to clarify my claims.
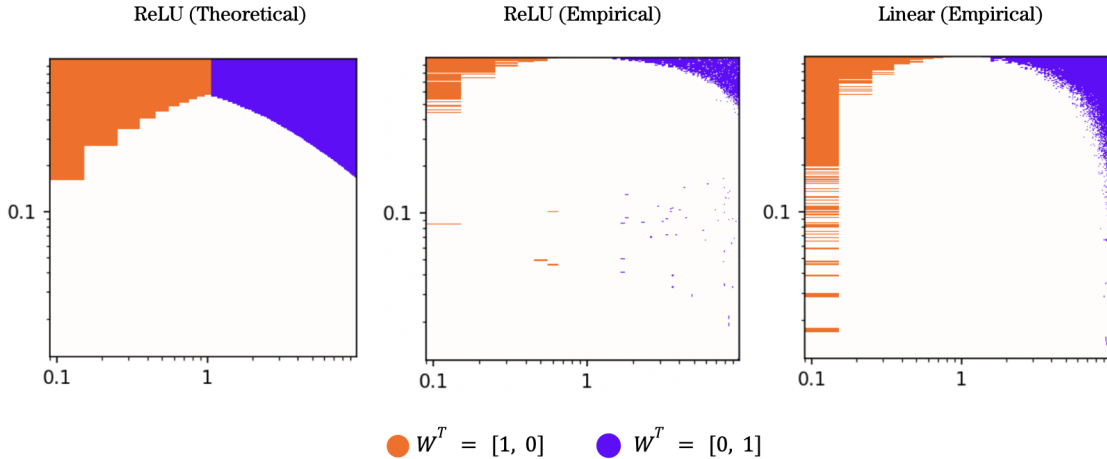

Figure 9: Theoretical and empirical superposition in one-neuron models. Unlike the authors of *Toy Models of Superposition*, I perform experiments on both models defined by $\text{ReLU}(W^T W x + b)$ and models without an activation function. Unlike the original authors however, my results have some extra limitations described in sections 4.2 and 4.3.

Note that the x-axis of the graphs in figure 9 represents the relative "importance" of the second output when calculating the loss. The y-axis represents the "feature density" of a model's training data (that is 1 - sparsity).

4.2 Replicating Theoretical Predictions

By using equations 2, 3 and 4, I was able to replicate the theoretical phase diagram from *Toy Models of Superposition*. The careful reader however, will notice that the theoretical prediction in figure 8 (graphic from *Toy Models of Superposition*) looks far less choppy than the one in figure 9 (graphic generated for this replication). This is because in the theoretical prediction I generated, steps between each value on the x-axis are much larger. This, however, is intentional and is beneficial for two reasons. First, by creating larger steps between numbers on the x-axis, I was able to intentionally tweak the step size to make the numbers easiest to work with when generating the graph. It is also beneficial because larger step sizes means that the graph includes far fewer theoretical models. This meant that attempting to construct an empirical version of the theoretical graph required training fewer real models. Because my current setup is very compute-constrained, this made performing the tasks in the following subsections much more manageable.
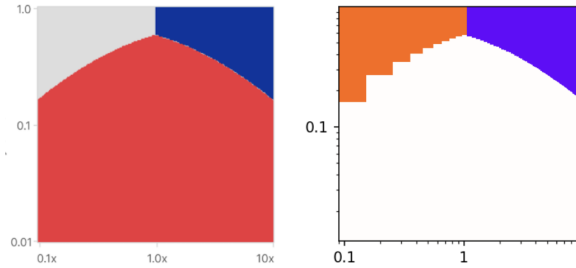

Figure 10: Comparison of theoretical predictions for internal structure of one-neuron models. Figure from *Toy Models of Superposition* is shown on the left while replication for this paper shown on right.

Important Point of Clarification: Although figure 9 includes a graph for both linear and ReLU empirical results, the theoretical predictions are strictly calculated assuming a ReLU activation function on the model. We discuss the linear findings in section 4.4, but the main focus of this replication is on the ReLU results.

4.3 Training 1,000,000 ReLU Models

Theoretical predictions in figure 9 account for 1000 different levels of feature sparsity and 100 levels of relative feature importance. As a result, in order to test the predictions empirically, it was necessary to train batches of 100,000 models. In order to create a good sample to detect trends, I trained 10 batches of 100,000 ReLU models, for a total of one million models across training. In my compute-constrained environment, this proved to be a significant challenge. For this reason, I had to make a few important considerations when writing code to train the ReLU models. I used similar methods to train a set of linear models but this is discussed separately in section 4.4.

The most obvious thing I did to improve performance when training the ReLU models was to not train each model independently. Instead, I constructed a PyTorch tensor which included all 100,000 models and used built-in PyTorch functions to train the models simultaneously. This helps PyTorch optimize training by leveraging features like parallel operations.

Another thing I did to make training more efficient was to initialize the weights and biases in a way that helps each model reduce its loss faster and avoid local minima. I found it to be highly beneficial to

initialize the weights using a normal distribution with a standard deviation of 1.5 instead of 1.0 (which perhaps would be a more natural choice in this kind of setup). This change makes sense because in theory, the model should "want" to make each item in its weight matrix either 1 or -1. Increasing the standard deviation puts more numbers around this range. Similarly, I found that it is advantageous to decrease the standard deviation of the bias parameters from 1 to 0.5. This makes intuitive sense because there is no case where the one-neuron models should "want" to have large bias values. It is worth noting that initializing the weights and biases in this way does not eliminate the need for the models to be trained using gradient descent. Figure 11 shows how even with these weight initializations, the overall loss of the trained models is still initially very high and quickly decreases as the models are trained. The goal of initializing the weights and biases with unconventional standard deviations is not to eliminate the need for training but to rather help the models achieve lower losses and train more efficiently.
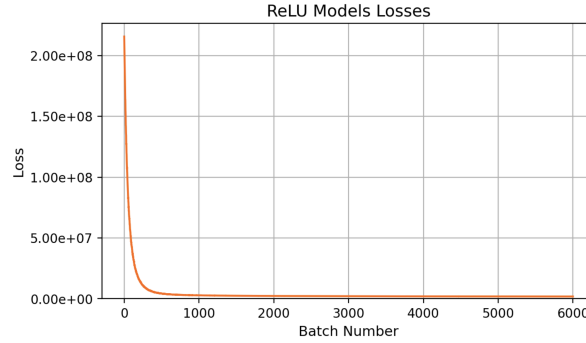

Figure 11: Loss for ReLU models trained in this replication. Note that each point on the graph is the loss for a batch of 64 examples of 100,000 models. The x-axis indicates the batch number in training.

The overall result of this training is shown below and compared with the findings from *Toy Models of Superposition*. Both diagrams show the empirical result of training one-neuron ReLU toy models.
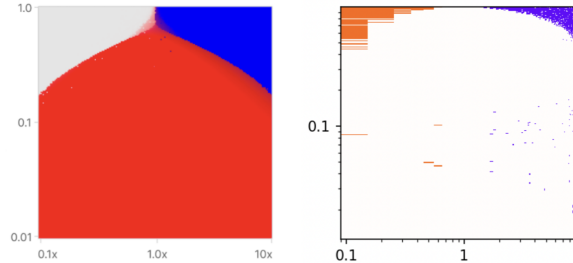

Figure 12: Empirical losses for one-neuron ReLU toy models. The graph on the left is taken from *Toy Models of Superposition*. The figure on the right was generated for this replication.

Although there are a few clear differences between the graphs in figure 12, the results are similar in the way that they each illustrate analogous overall trends. When a one-neuron ReLU model is trained with high feature density and with a loss that allocates a very low level of relative importance to the second output, the model tends to represent only the first input internally ($W \approx [1, 0]$). If the second feature has a high relative importance, only the second input is represented in W ($W \approx [0, 1]$). As the feature density of the model's training data decreases, the model has an incentive to not map features in this way but to rather represent them in superposition. This is the same phenomenon we observed in figure 2 from the introduction of this paper.

9

There are, however, also clearly a few important differences between the graphs in figure 12. The first is that the result generated by the authors of *Toy Models of Superposition* matches the theoretical expectation much better than the graph generated for this replication. This could be partially because of slightly different methods used in interpreting the results of training these one-neuron models. In the original paper, the authors train 10 batches of real models saying they "average over the results, discarding the model with the highest loss." I was not able to follow this exact process because it is not entirely clear what it means in this context to "average over the results." Because a model with $W = [1, 0]$ behaves in the exact same way as a model with $W = [-1, 0]$ it doesn't make sense to average over the models making $W_{avg} = [0, 0]$. Likewise, it does not make sense to take the absolute value of each weight matrix before averaging because in some cases having a negative number does make the model behave differently (e.g., $W = [1, -1]$). As a result, I took a slightly different approach. For each model I trained, I kept track of whether the weight matrix W represented only the first feature or only the second feature. If W represented neither or both, I skipped it. At the end of training, if more than half of the models trained at a given point mapped only the first feature, that point was colored orange on the resulting graph. If more than half the models mapped only the second feature, the point was colored purple. Using this method may have been more strict than the tactics used by the original authors making my graph more bare and less similar to the theoretical predictions.

Another factor that could have led to a discrepancy between the theoretical and empirical findings may have been the non-ideal environment in which one-neuron ReLU models were trained in. Because my setup was compute-constrained, I was not able to iterate and experiment as much as I wish I could have in training the models. As a result, I expect the loss I achieved could have been significantly lower in an ideal training environment.

4.4 Training 100,000 Linear Models

In addition to training a round of one-neuron ReLU models, I also trained a group without an activation function. Unlike the group of ReLU models in the previous section, I only trained each linear model once. This proved to be enough to get a general idea of how the models behave in practice.

In theory, these linear models should not exhibit superposition. It is in the best interest of these linear models to represent only one feature in the weight matrix—that is W should be $[1, 0]$ or $[0, 1]$. In practice, however, these linear models do not always do exactly what they "should" do. Although each model was trained extensively—6,000 batches of 128 examples—they did not always do what would limit the loss the most in theory. The outcome looked like figure 13 where the x-axis is once again relative feature importance of the second output and the y axis is feature density of the training data.
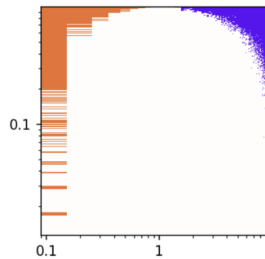

Figure 13: Empirical losses for one-neuron linear toy models.

This example shows something somewhat obvious: even if a model doesn't always do what is most advantageous for it, it is more likely to do the things it is incentivized to do. In theory, this graph should be all orange and purple. In practice, however, it only maps features like $[1, 0]$ or $[0, 1]$ when it is highly favorable. This illustrates a limitation of the research of *Toy Models of Superposition*: in practice it can be fairly complex to predict whether a model should represent a feature in superposition.

5 Conclusion

This replication demonstrates that it is possible for models to represent features in superposition. It also shows that the phenomenon of superposition is somewhat predictable. For example, models trained on sparse data are more likely to represent the features in superposition. On the other hand, if one feature is exceedingly important it is very unlikely that that feature will be represented in superposition. Overall, this replication shows that there are reasonable frameworks which machine learning researchers can apply to understand why a model represents information in the way it does.

With that being said, there are some serious limitations to thinking about neural networks in this way. In all the examples in this replication, the result was highly dependent on the exact training conditions such as learning rate and batch size. Thus, I believe that it is wise to be cautious when making broad claims about models such as "model x will represent less information in superposition than model y because model x is larger." The truth is that there are a number of important factors that influence whether or not a model will represent information orthogonally or in superposition.

References

[1] Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah. Toy models of superposition, 2022.